

# s-customer-segmentation-clustering

April 1, 2024

## 0.1 Project Name : Customer Segmentation Clustering

### 0.1.1 Contribution : Individual

The aim of this project is to perform clustering on the data based on the 'Occupation' column to identify group patterns that may exist between different occupations.

## 0.2 Data Dictionary

| Column Name     | Description                                     |
|-----------------|---|
| ID              | Shows a unique identifier of a customer         |
| Sex             | Gender of a customer. In this dataset there are |
| Marital Status  | Marital status of a customer                    |
| Age             | The age of a customer                           |
| Education       | Level education of a customer                   |
| Income          | Self-reported annual income of the customer     |
| Occupation      | Category of occupation of a customer            |
| Settlement Size | The size of the city that the customer lives in |

```
[1]: # import library
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
%matplotlib inline
sns.set(palette='pastel', style='whitegrid')
warnings.filterwarnings('ignore')

from sklearn.cluster import KMeans
from sklearn.metrics import confusion_matrix, classification_report, \
    silhouette_score

[2]: # load dataset
dataset = '/content/segmentation data.csv'

df = pd.read_csv(dataset)
```

```
df.head()
```

```
[2]:
```

|   | ID        | Sex | Marital status | Age | Education | Income | Occupation \ |
|---|-----------|-----|----------------|-----|-----------|--------|--------------|
| 0 | 100000001 | 0   | 0              | 67  | 2         | 124670 | 1            |
| 1 | 100000002 | 1   | 1              | 22  | 1         | 150773 | 1            |
| 2 | 100000003 | 0   | 0              | 49  | 1         | 89210  | 0            |
| 3 | 100000004 | 0   | 0              | 45  | 1         | 171565 | 1            |
| 4 | 100000005 | 0   | 0              | 53  | 1         | 149031 | 1            |

|   | Settlement size |
|---|-----------------|
| 0 | 2               |
| 1 | 2               |
| 2 | 0               |
| 3 | 1               |
| 4 | 1               |

sex : 0 female and 1 male

marital status : 0 single and 1 non single

### 0.3 Data Preprocessing Part 1

```
[3]: # drop the ID column because not needed
df.drop('ID', axis=1, inplace=True)
```

```
[4]: # checking the shape of dataset
df.shape
```

```
[4]: (2000, 7)
```

```
[5]: # checking detail information about dataset
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   Sex              2000 non-null   int64  
1   Marital status   2000 non-null   int64  
2   Age              2000 non-null   int64  
3   Education         2000 non-null   int64  
4   Income           2000 non-null   int64  
5   Occupation        2000 non-null   int64  
6   Settlement size   2000 non-null   int64  
dtypes: int64(7)
memory usage: 109.5 KB
```

```
[6]: # show the descriptive statistics of the dataset
df.describe()
```

```
[6]:
```

|       | Sex         | Marital status | Age         | Education   | Income \      |
|-------|-------------|----------------|-------------|-------------|---------------|
| count | 2000.000000 | 2000.000000    | 2000.000000 | 2000.000000 | 2000.000000   |
| mean  | 0.457000    | 0.496500       | 35.909000   | 1.03800     | 120954.419000 |
| std   | 0.498272    | 0.500113       | 11.719402   | 0.59978     | 38108.824679  |
| min   | 0.000000    | 0.000000       | 18.000000   | 0.00000     | 35832.000000  |
| 25%   | 0.000000    | 0.000000       | 27.000000   | 1.00000     | 97663.250000  |
| 50%   | 0.000000    | 0.000000       | 33.000000   | 1.00000     | 115548.500000 |
| 75%   | 1.000000    | 1.000000       | 42.000000   | 1.00000     | 138072.250000 |
| max   | 1.000000    | 1.000000       | 76.000000   | 3.00000     | 309364.000000 |

|       | Occupation  | Settlement size |
|-------|-------------|-----------------|
| count | 2000.000000 | 2000.000000     |
| mean  | 0.810500    | 0.739000        |
| std   | 0.638587    | 0.812533        |
| min   | 0.000000    | 0.000000        |
| 25%   | 0.000000    | 0.000000        |
| 50%   | 1.000000    | 1.000000        |
| 75%   | 1.000000    | 1.000000        |
| max   | 2.000000    | 2.000000        |

```
[7]: df['Education'].value_counts()
```

```
[7]: 1    1386
      2     291
      0     287
      3      36
      Name: Education, dtype: int64
```

```
0 : other
1 : high school
2 : university
3 : graduated school
```

```
[8]: df['Settlement size'].value_counts()
```

```
[8]: 0    989
      1    544
      2    467
      Name: Settlement size, dtype: int64
```

```
0 : small city
1 : mid city
```

2 : big city

```
[9]: df['Occupation'].value_counts()
```

```
[9]: 1    1113
     0     633
     2     254
     Name: Occupation, dtype: int64
```

0 : unemployed / unskilled

1 : skilled employee / official

2 : management / self-employed

```
[10]: # checking the duplicated value
      df.duplicated().sum()
```

```
[10]: 0
```

```
[11]: # checking the null value
      df.isnull().sum() * 100 / len(df)
```


```
[11]: Sex                0.0
      Marital status    0.0
      Age               0.0
      Education         0.0
      Income            0.0
      Occupation        0.0
      Settlement size   0.0
      dtype: float64
```

## 0.4 Eksploratory Data Analysis (EDA)

```
[12]: df.tail()
```

```
[12]:      Sex  Marital status  Age  Education  Income  Occupation  Settlement size
1995    1                0   47          1  123525            0                0
1996    1                1   27          1  117744            1                0
1997    0                0   31          0   86400            0                0
1998    1                1   24          1   97968            0                0
1999    0                0   25          0   68416            0                0
```


```
[13]: df['Sex'].value_counts().plot(kind='pie', autopct='%.2f%', figsize=(7,7))
      plt.title('Sex Distribution')
      plt.show()
```



Based on the pie plot distribution above, it can be seen that the representation of male sex is greater than that of female sex, with percentages of 54.30% and 45.70% respectively.


```
[14]: plt.figure(figsize=(15,5))
      ax = sns.histplot(data = df, x = 'Income')
      ax.set(title = 'Income Distribution')
      plt.vlines(df.Income.mean(),0, 190, color = 'red')
      plt.annotate('age mean = %.2f' % df.Income.mean(), (df.Income.mean(), 200))
```

```
[14]: Text(120954.419, 200, 'age mean = 120954.42')
```




```
[15]: plt.figure(figsize=(15,5))
      ax = sns.histplot(data = df, x = 'Age')
      ax.set(title = 'Age Distribution')
      plt.vlines(df.Age.mean(),0, 300, color = 'red')
      plt.annotate('age mean = %.2f years' % df.Age.mean(), (df.Age.mean(), 310))
```

```
[15]: Text(35.909, 310, 'age mean = 35.91 years')
```



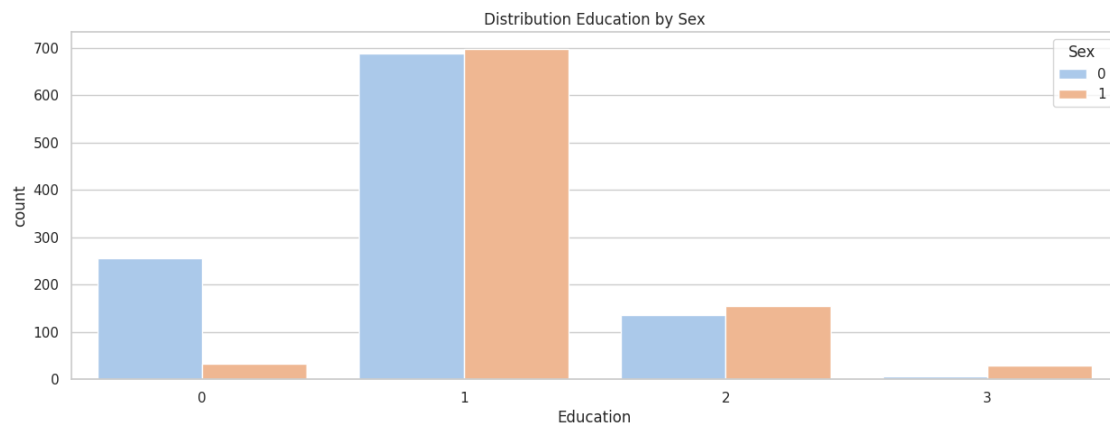
```
[16]: plt.figure(figsize=(15,5))
      sns.barplot(data=df, x='Education', y='Age',color='darkred').
      ↪set_title('Education Distribution')
```

```
[16]: Text(0.5, 1.0, 'Education Distribution')
```



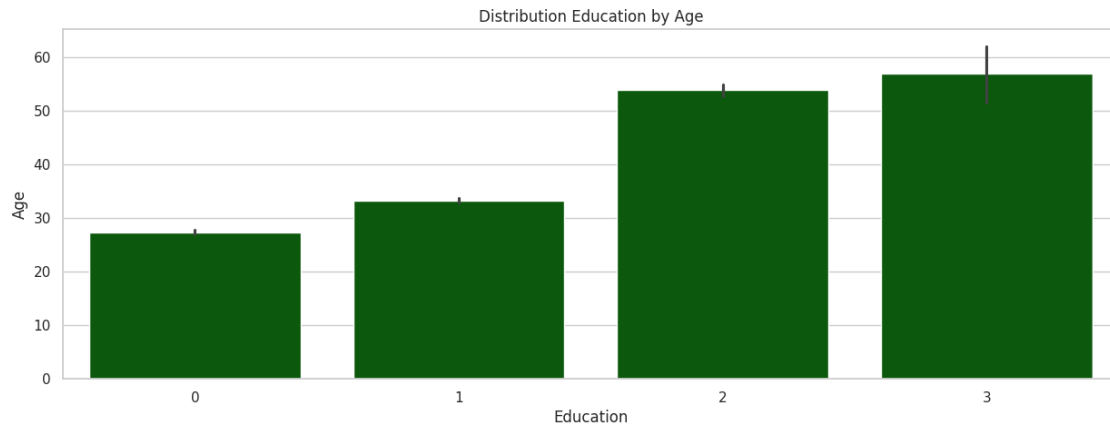
```
[17]: plt.figure(figsize=(15,5))
sns.countplot(data=df, x='Education', hue='Sex').set_title('Distribution_
↳ Education by Sex')
```

```
[17]: Text(0.5, 1.0, 'Distribution Education by Sex')
```



```
[18]: plt.figure(figsize=(15,5))
sns.barplot(data=df, x='Education', y='Age', color='darkgreen').
↳ set_title('Distribution Education by Age')
```


```
[18]: Text(0.5, 1.0, 'Distribution Education by Age')
```



```
[19]: plt.figure(figsize=(8,8))  
sns.heatmap(df.corr(), annot=True, cmap='BuPu').set_title('Correlation Plot')
```

```
[19]: Text(0.5, 1.0, 'Correlation Plot')
```






## 0.5 Data Preprocessing Part 2

```
[20]: num_clusters = range(1, 11)
      inertia = []
      for k in num_clusters:
          kmeans = KMeans(n_clusters=k)
          kmeans.fit(df)
          inertia.append(kmeans.inertia_)
```

```
[21]: plt.figure(figsize=(8, 6))
      plt.plot(num_clusters, inertia, marker='o')
```

```
plt.title('The elbow method')
plt.xticks(num_clusters)
plt.show()
```




According to the elbow method we can assume 3 clusters for our data

```
[22]: SEED = np.random.seed(300)
      kmeans = KMeans(n_clusters=4, random_state=SEED)
      labels = kmeans.fit_predict(df)
      df['cluster'] = labels
```

```
[23]: sns.scatterplot(data = df, x = 'Age', y = 'Income', hue = 'cluster',
                      palette='colorblind')
```

```
[23]: <Axes: xlabel='Age', ylabel='Income'>
```



```
[24]: silhouette_avg = silhouette_score(df, df.cluster)

print(f"Silhouette Score: {silhouette_avg}")
```

Silhouette Score: 0.5428478923817803

The silhouette score ranges from -1 to 1.

A higher silhouette score indicates better-defined clusters. Evaluate the score based on the following guidelines:

0.71 - 1.0: Excellent clustering.

0.51 - 0.70: Reasonable clustering.

0.26 - 0.50: Poor clustering.

Less than 0.25: Very poor clustering

Let's try different numbers of clusters, i will choose 3 and 5 cluster therefore it is closest to 4

```
[25]: # cluster 3
SEED = np.random.seed(300)
df.drop(columns = 'cluster')
kmeans = KMeans(n_clusters = 3, random_state = SEED)
```

```

labels = kmeans.fit_predict(df)
df['cluster'] = labels
silhouette_avg = silhouette_score(df, df.cluster)
print(f"Silhouette Score: {silhouette_avg}")

```

Silhouette Score: 0.5124711688681454

```

[26]: # cluster 5
SEED = np.random.seed(300)
df.drop(columns = 'cluster')
kmeans = KMeans(n_clusters = 5, random_state = SEED)
labels = kmeans.fit_predict(df)
df['cluster'] = labels
silhouette_avg = silhouette_score(df, df.cluster)
print(f"Silhouette Score: {silhouette_avg}")

```

Silhouette Score: 0.5401417577104871

We can conclude that the best number of clusters was 4 according to silhouette score.

```

[27]: # Getting data with 4 clusters
SEED = np.random.seed(300)
df.drop(columns = 'cluster')
kmeans = KMeans(n_clusters = 4, random_state = SEED)
labels = kmeans.fit_predict(df)
df['cluster'] = labels
silhouette_avg = silhouette_score(df, df.cluster)
print(f"Silhouette Score: {silhouette_avg}")

```

Silhouette Score: 0.5428478923817803

```


[28]: sns.scatterplot(data = df, x = 'Age', y = 'Income', hue = 'cluster',
    ↪ palette='colorblind')

```

```

[28]: <Axes: xlabel='Age', ylabel='Income'>

```



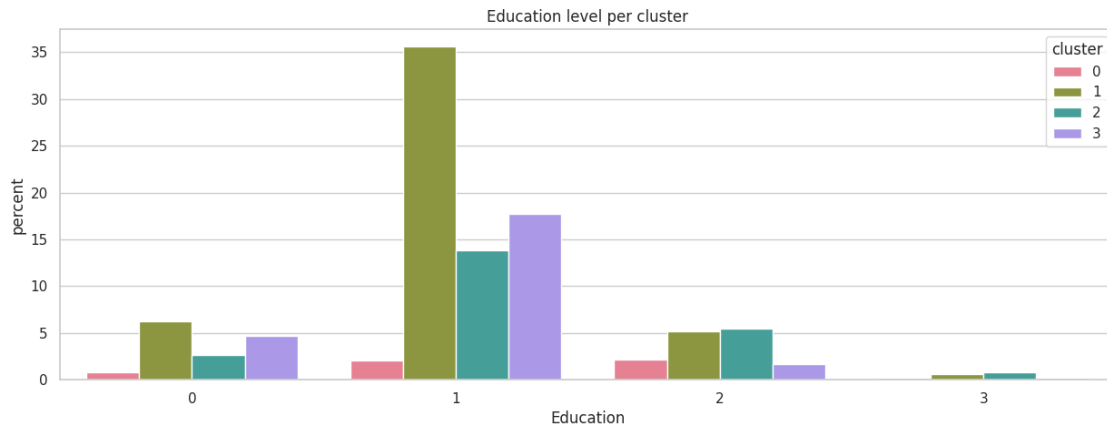
```
[29]: # Which Cluster has the best income?
plt.figure(figsize=(15,5))
sns.barplot(df, x = 'cluster', y = 'Income').set_title('Income per cluster')
```

```
[29]: Text(0.5, 1.0, 'Income per cluster')
```



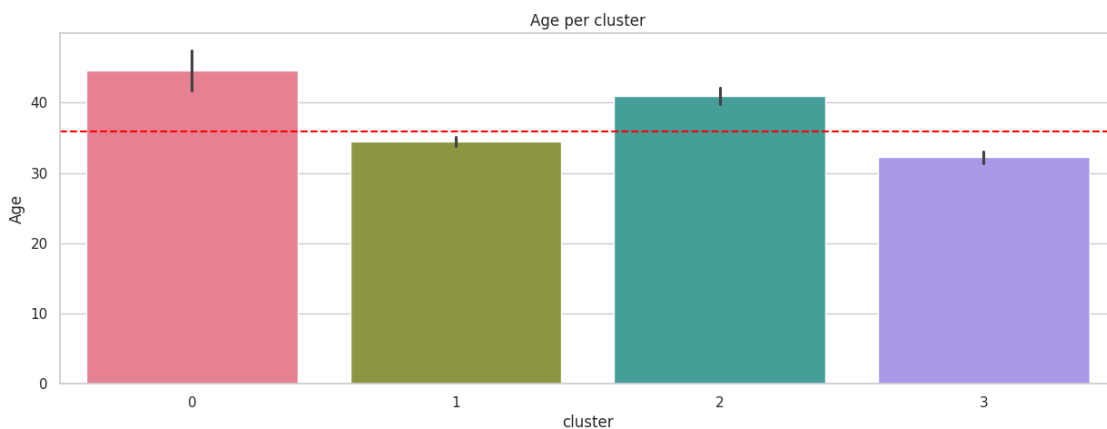
```
[30]: # Which cluster has the best education level?
plt.figure(figsize=(15,5))
sns.countplot(df, x = 'Education', hue = 'cluster',stat='percent', palette = 'husl',
↪sns.color_palette('husl', 4)).set_title('Education level per cluster')
```

```
[30]: Text(0.5, 1.0, 'Education level per cluster')
```



```
[31]: # What is the average age per cluster?
plt.figure(figsize=(15,5))
sns.barplot(df, x = 'cluster', y = 'Age', palette = sns.color_palette('husl', 4))
↪).set_title('Age per cluster')
plt.axhline(df['Age'].mean(), 0,4, linestyle = '--', color = 'red')
```

```
[31]: <matplotlib.lines.Line2D at 0x7c5679a8ac80>
```



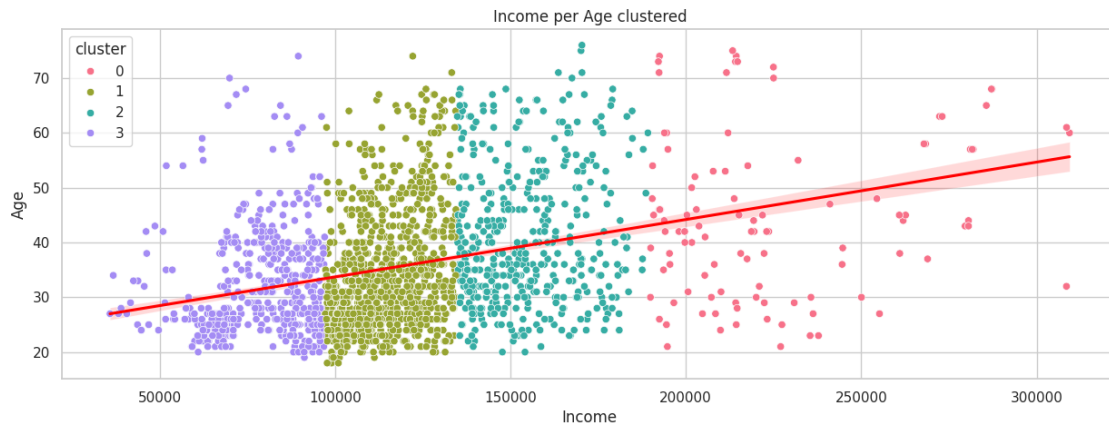
```
[32]: # Is Salary correlated with Age?
plt.figure(figsize=(15,5))
```

```

sns.regplot(df, x = 'Income', y= 'Age',scatter_kws={'alpha':0.0},
            line_kws={'color':'red'})
sns.scatterplot(df, x='Income',y='Age',hue='cluster', palette = sns.
               color_palette('husl', 4))
plt.title('Income per Age clustered')

```

[32]: Text(0.5, 1.0, 'Income per Age clustered')



We can see that elder people earns more money than yonger people and we have it clusterized

### 0.5.1 Conclusion

From the results of the k means clustering that we have done, it can be concluded that there are 4 clusters for customer segmentation, starting from income analysis per cluster, education per cluster and there are still many analyzes that can be done from these 4 clusters to get a better analysis.