

predicting-survival-titanic

April 9, 2024

0.1 Project Name : Predicting Survival on Titanic using Artificial Neural Networks

0.1.1 Contribution : Individual

0.2 Data Dictionary

Column Name	Description
Pclass	Ticket class indicating the socio-economic status of the passenger
Survived	A binary indicator that shows whether the passenger survived (1) or not (0)
Name	The full name of the passenger
Sex	The gender of the passenger, denoted as either male or female
Age	The age of the passenger in years
SibSp	The number of siblings or spouses aboard the Titanic for the respective passenger
Parch	The number of parents or children aboard the Titanic for the respective passenger
Ticket	The ticket number assigned to the passenger
Fare	The fare paid by the passenger for the ticket
Cabin	The cabin number assigned to the passenger, if available
Embarked	The port of embarkation for the passenger
Boat	this column contains the identifier of the lifeboat they were rescued in
Body	this column contains the identification number of their recovered body, if applicable
Home.dest	The destination or place of residence of the passenger

```
[35]: # import library
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
sns.set(style='darkgrid', palette='Set2')
```

```
[36]: # load dataset
dataset = '/content/Titanic Dataset.csv'

df = pd.read_csv(dataset)
df.head()
```

```
[36]:
```

	pclass	survived	name	sex	\
0	1	1	Allen, Miss. Elisabeth Walton	female	
1	1	1	Allison, Master. Hudson Trevor	male	
2	1	0	Allison, Miss. Helen Loraine	female	
3	1	0	Allison, Mr. Hudson Joshua Creighton	male	
4	1	0	Allison, Mrs. Hudson J C (Bessie Waldo Daniels)	female	

	age	sibsp	parch	ticket	fare	cabin	embarked	boat	body	\
0	29.00	0	0	24160	211.3375	B5	S	2	NaN	
1	0.92	1	2	113781	151.5500	C22 C26	S	11	NaN	
2	2.00	1	2	113781	151.5500	C22 C26	S	NaN	NaN	
3	30.00	1	2	113781	151.5500	C22 C26	S	NaN	135.0	
4	25.00	1	2	113781	151.5500	C22 C26	S	NaN	NaN	

	home.dest
0	St Louis, MO
1	Montreal, PQ / Chesterville, ON
2	Montreal, PQ / Chesterville, ON
3	Montreal, PQ / Chesterville, ON
4	Montreal, PQ / Chesterville, ON

0.3 Data Preprocessing 1

```
[37]: # checking the shape dataset
df.shape
```

```
[37]: (1309, 14)
```

Dropping the unnecessary columns - name, cabin, boat, body, home.dest

```
[38]: # drop columns
drop_col = ['name', 'cabin', 'boat', 'body', 'home.dest']

df.drop(drop_col, axis=1, inplace=True)
df.dtypes
```

```
[38]: pclass      int64
survived      int64
sex           object
age          float64
sibsp         int64
```

```
parch          int64
ticket         object
fare          float64
embarked       object
dtype: object
```

```
[39]: # moved the 'survived' column to the last column
df = pd.concat([df.drop(columns=['survived']), df[['survived']]], axis=1)
df.sample(5)
```

```
[39]:
```

	pclass	sex	age	sibsp	parch	ticket	fare	embarked	\
375	2	male	28.0	0	0	248740	13.00	S	
973	3	male	NaN	0	0	S.O./P.P. 251	7.55	S	
583	2	female	40.0	0	0	C.A. 33595	15.75	S	
484	2	female	34.0	0	0	C.A. 34260	10.50	S	
1073	3	male	NaN	0	0	371060	7.75	Q	

	survived
375	0
973	0
583	1
484	1
1073	0

```
[40]: # shwoing 10 values for ticket column
df.ticket.values[:10]
```

```
[40]: array(['24160', '113781', '113781', '113781', '113781', '19952', '13502',
        '112050', '11769', 'PC 17609'], dtype=object)
```

because the ticket column is of the object data type, I will convert it to numeric

```
[41]: # convert ticket column to float
df.ticket = pd.to_numeric(df.ticket, errors='coerce')
```

```
[42]: # checking the dtype for ticket column
df.ticket.dtypes
```

```
[42]: dtype('float64')
```

```
[43]: # checking null value
df.isnull().sum()
```

```
[43]: pclass      0
sex           0
age          263
sibsp        0
```

```
parch      0
ticket     352
fare       1
embarked   2
survived   0
dtype: int64
```

```
[44]: # checking null value for age column
df[df.age.isnull()]
```

```
[44]:      pclass    sex  age  sibsp  parch  ticket    fare embarked  survived
15         1   male  NaN     0     0      NaN   25.9250         S         0
37         1   male  NaN     0     0  111427.0   26.5500         S         1
40         1   male  NaN     0     0  112379.0   39.6000         C         0
46         1   male  NaN     0     0  113798.0   31.0000         S         0
59         1  female  NaN     0     0   17770.0   27.7208         C         1
...
1293        3   male  NaN     0     0      NaN    8.0500         S         0
1297        3   male  NaN     0     0      NaN    7.2500         S         0
1302        3   male  NaN     0     0   2647.0    7.2250         C         0
1303        3   male  NaN     0     0   2627.0   14.4583         C         0
1305        3  female  NaN     1     0   2665.0   14.4542         C         0
```

[263 rows x 9 columns]

```
[45]: # Replace the null value in the 'age' column with the mean
df.age.fillna(df.age.mean(), inplace=True)
df.age.isnull().sum()
```

```
[45]: 0
```

```
[46]: # drop null value for ticket column
df.dropna(subset=['ticket'], inplace=True)
df.ticket.isnull().sum()
```

```
[46]: 0
```

```
[47]: # checking null value for fare and embarked column
df[(df.fare.isnull()) | (df.embarked.isnull())]
```

```
[47]:      pclass    sex  age  sibsp  parch  ticket    fare embarked  survived
168         1  female  38.0     0     0  113572.0   80.0         NaN         1
284         1  female  62.0     0     0  113572.0   80.0         NaN         1
1225        3   male  60.5     0     0   3701.0    NaN         S         0
```

```
[48]: # Replace the null value in the 'fare' column with the mean
df.fare.fillna(df.fare.mean(), inplace=True)
```

```
# drop null value for embarked column
df.dropna(subset=['embarked'], inplace=True)
```

```
[49]: # checking null value again
df.isnull().sum()
```

```
[49]: pclass      0
      sex        0
      age        0
      sibsp      0
      parch      0
      ticket     0
      fare       0
      embarked   0
      survived   0
      dtype: int64
```

Because all columns no longer have null values, will be do checking duplicated values

```
[50]: # checking duplicated values
df[df.duplicated()]
```

```
[50]:
```

	pclass	sex	age	sibsp	parch	ticket	fare	embarked	\
384	2	male	29.881138	0	0	239853.0	0.0000	S	
438	2	female	24.000000	1	2	220845.0	65.0000	S	
528	2	male	29.881138	0	0	239853.0	0.0000	S	
658	3	female	0.750000	2	1	2666.0	19.2583	C	
714	3	male	32.000000	0	0	1601.0	56.4958	S	
858	3	male	29.881138	0	0	1601.0	56.4958	S	
945	3	male	29.881138	0	0	1601.0	56.4958	S	
956	3	female	29.881138	3	1	4133.0	25.4667	S	
957	3	female	29.881138	3	1	4133.0	25.4667	S	
1002	3	female	29.881138	2	0	367226.0	23.2500	Q	
1035	3	male	29.881138	1	1	2661.0	15.2458	C	
1043	3	female	29.881138	1	0	367230.0	15.5000	Q	
1185	3	male	29.881138	2	0	2662.0	21.6792	C	
1186	3	male	29.881138	2	0	2662.0	21.6792	C	

	survived
384	0
438	1
528	0
658	1
714	1
858	1
945	1

956	0
957	0
1002	1
1035	1
1043	1
1185	0
1186	0

In my opinion, it is not necessary to delete duplicate values, so we continue to the next step

```
[51]: # descriptive statistics
df.describe(include='all')
```

```
[51]:
```


	pclass	sex	age	sibsp	parch	ticket \
count	955.000000	955	955.000000	955.000000	955.000000	9.550000e+02
unique	NaN	2	NaN	NaN	NaN	NaN
top	NaN	male	NaN	NaN	NaN	NaN
freq	NaN	615	NaN	NaN	NaN	NaN
mean	2.342408	NaN	29.657154	0.450262	0.341361	2.493228e+05
std	0.831150	NaN	12.765045	0.840776	0.765282	4.431056e+05
min	1.000000	NaN	0.330000	0.000000	0.000000	6.800000e+02
25%	2.000000	NaN	22.000000	0.000000	0.000000	1.995000e+04
50%	3.000000	NaN	29.881138	0.000000	0.000000	2.346860e+05
75%	3.000000	NaN	34.000000	1.000000	0.000000	3.474685e+05
max	3.000000	NaN	80.000000	4.000000	5.000000	3.101298e+06

	fare	embarked	survived
count	955.000000	955	955.000000
unique	NaN	3	NaN
top	NaN	S	NaN
freq	NaN	662	NaN
mean	27.701066	NaN	0.378010
std	37.832230	NaN	0.485144
min	0.000000	NaN	0.000000
25%	7.883350	NaN	0.000000
50%	13.000000	NaN	0.000000
75%	27.900000	NaN	1.000000
max	263.000000	NaN	1.000000

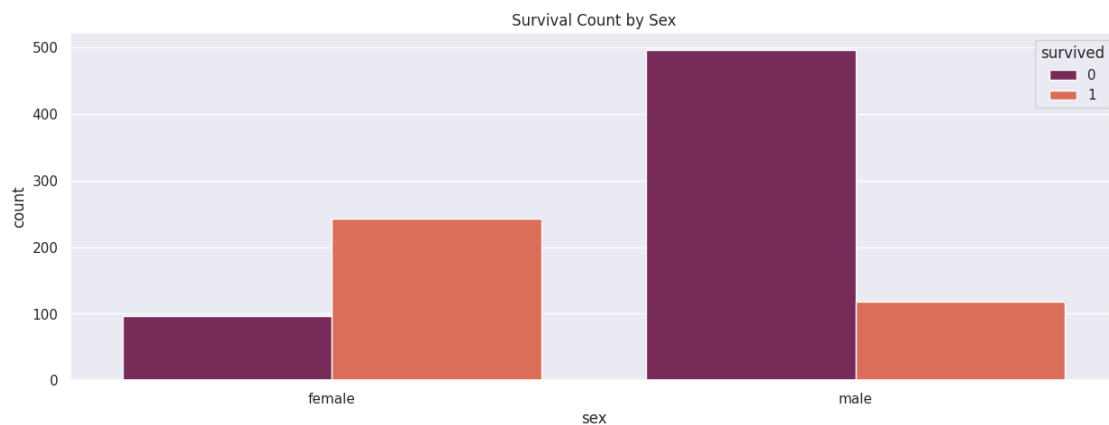
0.4 Eksploratory Data Analysis (EDA)

```
[53]: plt.figure(figsize=(6,6))
df['survived'].value_counts().plot(kind='pie', autopct='%.2f%%', labels=['no', 'yes'])
plt.title('Survived vs Not Survived Distribution')
plt.show()
```


Survived vs Not Survived Distribution



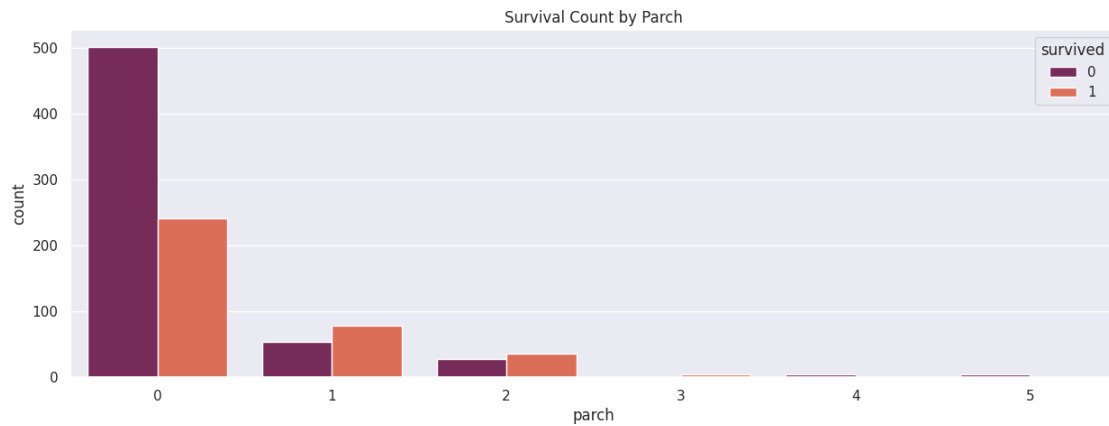
```
[54]: plt.figure(figsize=(15,5))
sns.countplot(data=df, x='sex', hue='survived', palette='rocket')
plt.title('Survival Count by Sex')
plt.show()
```




```
[55]: plt.figure(figsize=(15,5))
sns.countplot(data=df, x='pclass', hue='survived', palette='rocket')
plt.title('Survival Count by Pclass')
plt.show()
```



```
[56]: plt.figure(figsize=(15,5))
sns.countplot(data=df, x='parch', hue='survived', palette='rocket')
plt.title('Survival Count by Parch')
plt.show()
```



```
[57]: plt.figure(figsize=(15,5))
sns.countplot(data=df, x='embarked', hue='survived', palette='rocket')
plt.title('Survival Count by Embarked')
plt.show()
```

```
[58]: plt.figure(figsize=(15, 5))
sns.histplot(data=df, x='age', hue='survived', multiple='stack', kde=True)
plt.title('Survived Distribution by Age')
plt.xlim(0,100)
plt.show()
```



0.5 Data Preprocessing Part 2

```
[59]: # filter columns that have the object data type
for col in df:
    if df[col].dtypes == 'object':
        print(f'{col} : {df[col].unique()}')
```

```
sex : ['female' 'male']
embarked : ['S' 'C' 'Q']
```

```
[60]: # label encoding
from sklearn.preprocessing import LabelEncoder

var = ['sex', 'embarked']
le = LabelEncoder()

for i in var:
    le.fit(df[i].unique())
    df[i]=le.transform(df[i])
    print(i,df[i].unique())
```

```
sex [0 1]
embarked [2 0 1]
```

```
[61]: df.sample(5)
```

```
[61]:
```


	pclass	sex	age	sibsp	parch	ticket	fare	embarked	\
589	2	0	29.000000	0	2	29103.0	23.0000	2	
469	2	0	29.881138	0	0	226593.0	12.3500	1	
440	2	0	48.000000	1	2	220845.0	65.0000	2	
916	3	0	4.000000	0	1	349256.0	13.4167	0	
128	1	1	47.000000	0	0	111320.0	38.5000	2	

	survived
589	1
469	1
440	1
916	1
128	0

```
[62]: # normalize the continuous variables
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()

df[['ticket', 'fare']] = scaler.fit_transform(df[['ticket', 'fare']])
```

```
[63]: # correlation matrix
plt.figure(figsize=(10,10))
sns.heatmap(df.corr(),annot=True,cmap='BuPu')
plt.title('Correlation Matrix')
plt.show()
```



```
[64]: X = df.drop('survived', axis=1)
      y = df['survived']
```

```
[65]: # train test split
      from sklearn.model_selection import train_test_split

      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪ random_state=42)
```

```
[66]: len(X_train.columns)
```

```
[66]: 8
```

to predict whether the titanic people survived or not i will use the ANN model

```
[71]: import tensorflow as tf
      from tensorflow import keras

      # define the ANN model
      model = keras.Sequential([
          keras.layers.Dense(64, input_shape=(8,), activation='relu'),
          keras.layers.Dense(32, activation='relu'),
          keras.layers.Dense(1, activation='sigmoid')
      ])

      # compile the model
      model.compile(
          optimizer='adam',
          loss='binary_crossentropy',
          metrics=['accuracy']
      )

      # fitting the model
      model.fit(X_train, y_train, epochs=10)
```

Epoch 1/10

24/24 [=====] - 1s 4ms/step - loss: 1.5803 - accuracy: 0.5157

Epoch 2/10

24/24 [=====] - 0s 2ms/step - loss: 0.6895 - accuracy: 0.5929

Epoch 3/10

24/24 [=====] - 0s 2ms/step - loss: 0.6286 - accuracy: 0.6453

Epoch 4/10

24/24 [=====] - 0s 2ms/step - loss: 0.6031 - accuracy: 0.6662

Epoch 5/10

24/24 [=====] - 0s 2ms/step - loss: 0.5886 - accuracy: 0.6859

Epoch 6/10

24/24 [=====] - 0s 2ms/step - loss: 0.5694 - accuracy: 0.6963

Epoch 7/10

24/24 [=====] - 0s 2ms/step - loss: 0.5642 - accuracy: 0.7055

Epoch 8/10

```
24/24 [=====] - 0s 3ms/step - loss: 0.5504 - accuracy: 0.7160
Epoch 9/10
24/24 [=====] - 0s 2ms/step - loss: 0.5396 - accuracy: 0.7474
Epoch 10/10
24/24 [=====] - 0s 2ms/step - loss: 0.5332 - accuracy: 0.7421
```

[71]: <keras.src.callbacks.History at 0x7e0acc469c30>

```
[90]: # evaluate the model
model.evaluate(X_test, y_test)
```

```
6/6 [=====] - 0s 6ms/step - loss: 0.5317 - accuracy: 0.7801
```

[90]: [0.5317214131355286, 0.7801046967506409]

```
[91]: # do a predict and show 5 result the predict
yp = model.predict(X_test)
yp[:5]
```

```
6/6 [=====] - 0s 2ms/step
```

[91]: array([[0.59424615],
[0.7199342],
[0.21613458],
[0.4685506],
[0.22969034]], dtype=float32)

```
[92]: # filtering where val < 0.5 convert to 0 and else convert to 1
y_pred = []
for i in yp:
    if i < 0.5:
        y_pred.append(0)
    else:
        y_pred.append(1)
```

```
[93]: # show 5 pred result
y_pred[:5]
```

[93]: [1, 1, 0, 0, 0]

```
[94]: # show 5 the test result
y_test[:5]
```

```
[94]: 293      1
      626      0
      976      0
      1044     1
      571      0
      Name: survived, dtype: int64
```

of the 5 samples we took, it can be seen that only 3 predictions were correct for the 5 samples taken

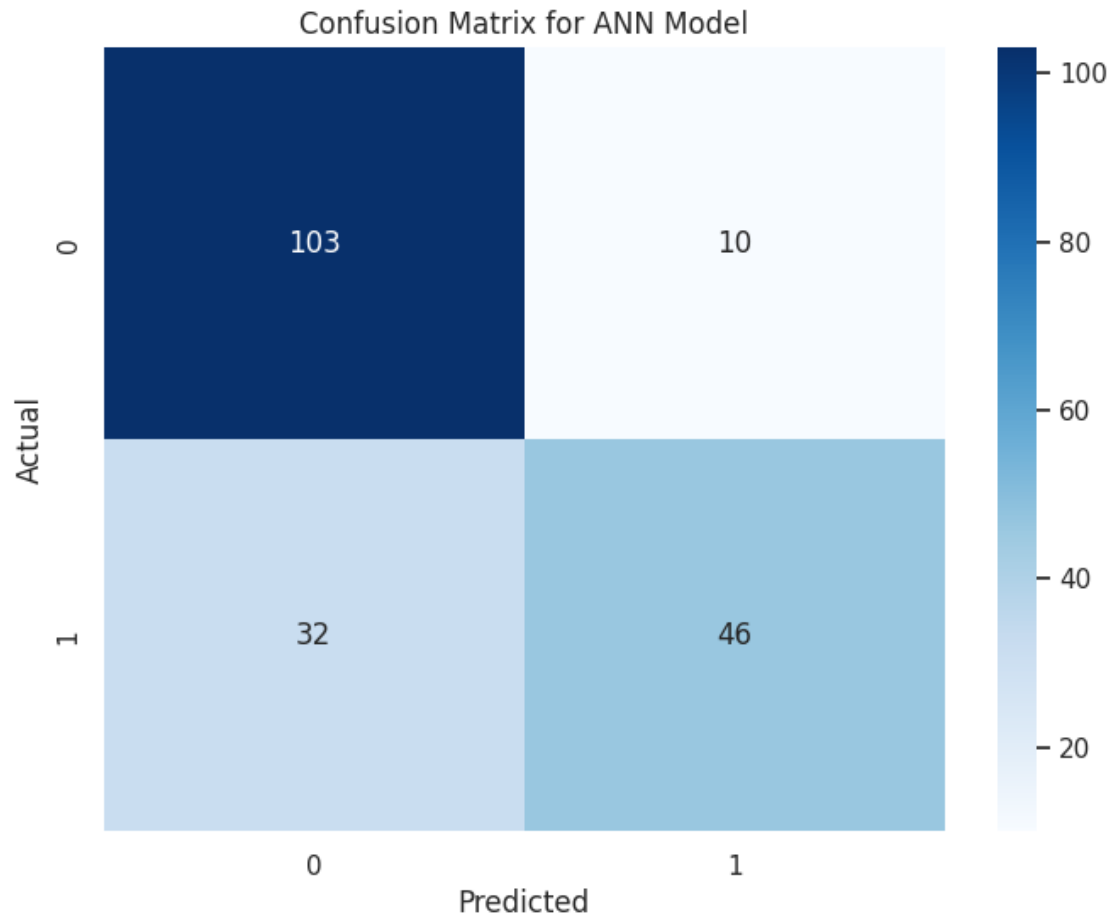
0.6 Model Evaluation

```
[84]: # classification report
      from sklearn.metrics import confusion_matrix, classification_report

      print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.76	0.91	0.83	113
1	0.82	0.59	0.69	78
accuracy			0.78	191
macro avg	0.79	0.75	0.76	191
weighted avg	0.79	0.78	0.77	191

```
[83]: # confusion matrix heatmap
      plt.figure(figsize=(8,6))
      sns.heatmap(confusion_matrix(y_test,y_pred),annot=True,fmt='d',cmap='Blues')
      plt.xlabel('Predicted')
      plt.ylabel('Actual')
      plt.title('Confusion Matrix for ANN Model')
      plt.show()
```



```
[95]: # distribution plot
ax = sns.distplot(y_test, hist=False, color='r', label='Actual Value')
sns.distplot(y_pred, hist=False, color='b', label='Predicted Values' , ax=ax)
plt.legend()
```

```
[95]: <matplotlib.legend.Legend at 0x7e0abfe789d0>
```

